

Programming with Python

(index.html)

Discussion

Rules of Debugging

1. Fail early, fail often.
2. Always initialize from data.
3. Know what it's supposed to do.
4. Make it fail every time.
5. Make it fail fast.
6. Change one thing at a time, for a reason.
7. Keep track of what we've done.
8. Be humble.
9. Test the simple things first.


And remember, a week of hard work can sometimes save you an hour of thought.

The Call Stack

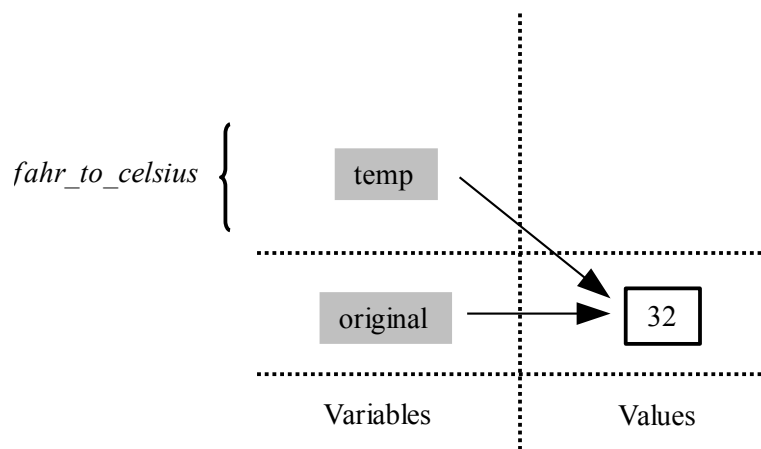
Let's take a closer look at what happens when we call `fahr_to_celsius(32.0)`. To make things clearer, we'll start by putting the initial value 32.0 in a variable and store the final result in one as well:

```
original = 32.0
final = fahr_to_celsius(original)
```

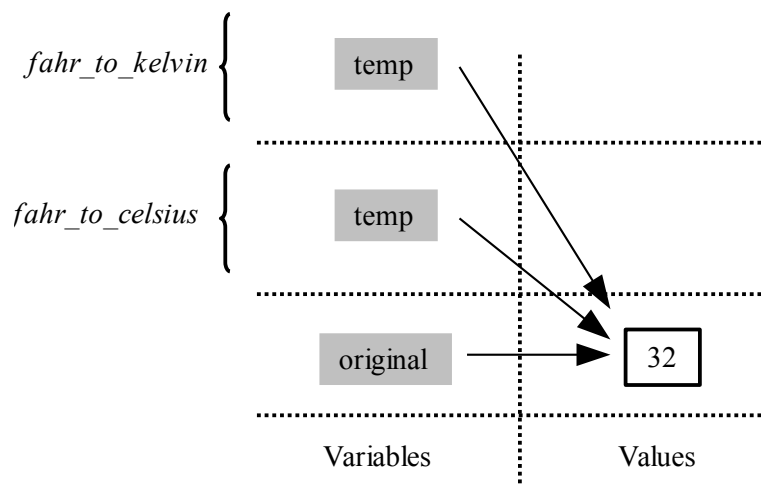
The diagram below shows what memory looks like after the first line has been executed:



When we call `fahr_to_celsius`, Python doesn't create the variable `temp` right away. Instead, it creates something called a [stack frame](#) (reference.html#stack-frame) to keep track of the variables defined by `fahr_to_kelvin`. Initially, this stack frame only holds the value of `temp`:

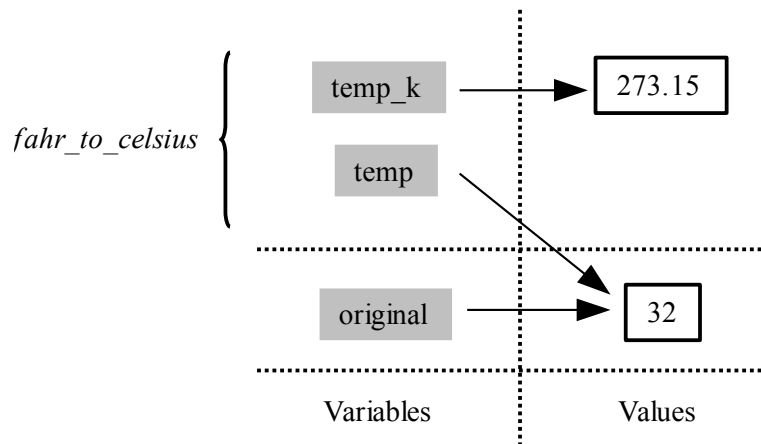


When we call `fahr_to_kelvin` inside `fahr_to_celsius`, Python creates another stack frame to hold `fahr_to_kelvin`'s variables:

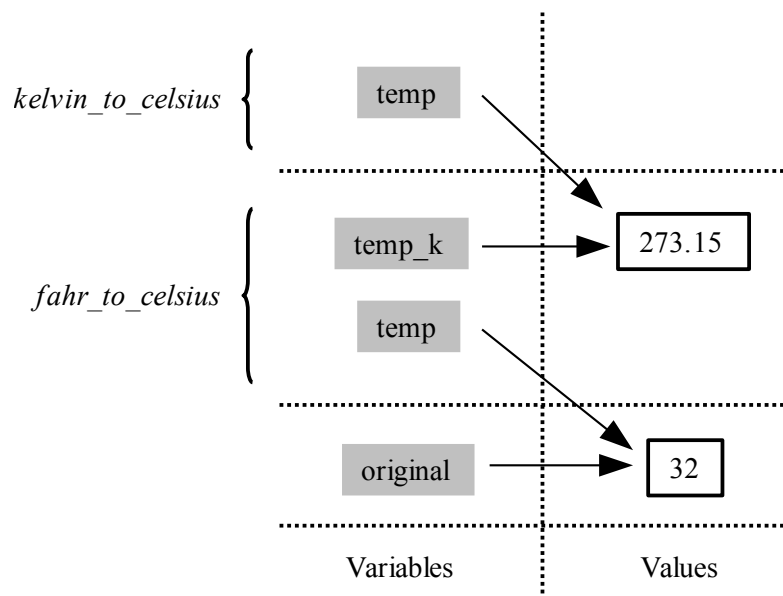


It does this because there are now two variables in play called `temp`: the parameter to `fahr_to_celsius`, and the parameter to `fahr_to_kelvin`. Having two variables with the same name in the same part of the program would be ambiguous, so Python (and every other modern programming language) creates a new stack frame for each function call to keep that function's variables separate from those defined by other functions.

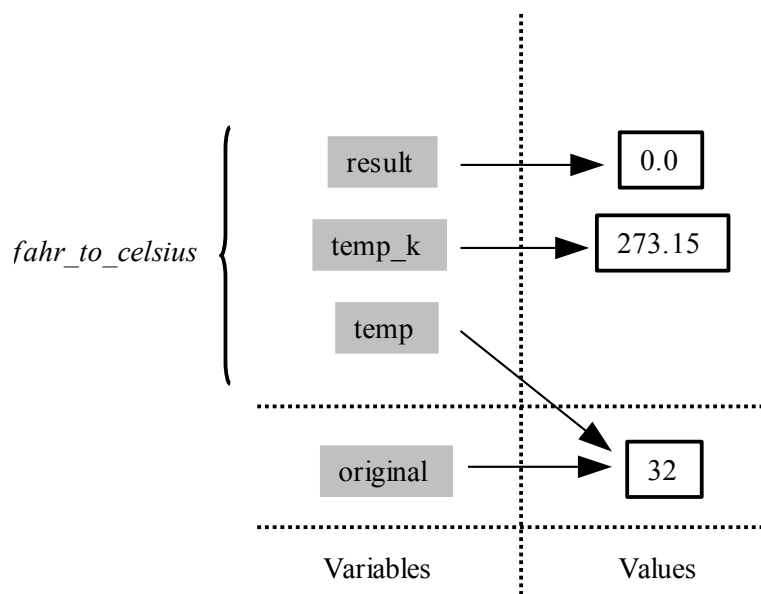
When the call to `fahr_to_kelvin` returns a value, Python throws away `fahr_to_kelvin`'s stack frame and creates a new variable in the stack frame for `fahr_to_celsius` to hold the temperature in Kelvin:




It then calls `kelvin_to_celsius`, which means it creates a stack frame to hold that function's variables:



Once again, Python throws away that stack frame when `kelvin_to_celsius` is done and creates the variable `result` in the stack frame for `fahr_to_celsius`:



Finally, when `fahr_to_celsius` is done, Python throws away its stack frame and puts its result in a new variable called `final` that lives in the stack frame we started with:



This final stack frame is always there; it holds the variables we defined outside the functions in our code. What it doesn't hold is the variables that were in the various stack frames. If we try to get the value of `temp` after our functions have finished running, Python tells us that there's no such thing:

```
print('final value of temp after all function calls:', temp)
```

```
-----
NameError                                Traceback (most recent call last)
<ipython-input-12-ffd9b4dbd5f1> in <module>()
----> 1 print('final value of temp after all function calls:', temp)

NameError: name 'temp' is not defined
```

```
final value of temp after all function calls:
```

Why go to all this trouble? Well, here's a function called `span` that calculates the difference between the minimum and maximum values in an array:

```
import numpy

def span(a):
    diff = a.max() - a.min()
    return diff

data = numpy.loadtxt(fname='inflammation-01.csv', delimiter=',')
print('span of data', span(data))
```

```
span of data 20.0
```

Notice that `span` assigns a value to a variable called `diff`. We might very well use a variable with the same name to hold data:

```
diff = numpy.loadtxt(fname='inflammation-01.csv', delimiter=',')
print('span of data:', span(diff))
```

```
span of data: 20.0
```

We don't expect `diff` to have the value 20.0 after this function call, so the name `diff` cannot refer to the same thing inside `span` as it does in the main body of our program. And yes, we could probably choose a different name than `diff` in our main program in this case, but we don't want to have to read every line of NumPy to see what variable names its functions use before calling any of those functions, just in case they change the values of our variables.

The big idea here is [encapsulation](#) ([reference.html#encapsulation](#)), and it's the key to writing correct, comprehensible programs. A function's job is to turn several operations into one so that we can think about a single function call instead of a dozen or a hundred statements each time we want to do something. That only works if functions don't interfere with each other; if they do, we have to pay attention to the details once again, which quickly overloads our short-term memory.

Following the call stack

We previously wrote functions called `fence` and `outer`. Draw a diagram showing how the call stack changes when we run the following:

```
print(outer(fence('carbon', '+')))
```

Image Grids


Let's start by creating some simple heat maps of our own using a library called `ipythonblocks`. The first step is to create our own "image":

```
from ipythonblocks import ImageGrid
```

Unlike the `import` statements we have seen earlier, this one doesn't load the entire `ipythonblocks` library. Instead, it just loads `ImageGrid` from that library, since that's the only thing we need (for now).

Once we have `ImageGrid` loaded, we can use it to create a very simple grid of colored cells:

```
grid = ImageGrid(5, 3)
grid.show()
```



Just like a NumPy array, an `ImageGrid` has some properties that hold information about it:

```
print('grid width:', grid.width)
print('grid height:', grid.height)
print('grid lines on:', grid.lines_on)
```

```
grid width: 5
grid height: 3
grid lines on: True
```

The obvious thing to do with a grid like this is color in its cells, but in order to do that, we need to know how computers represent color. The most common schemes are [RGB](#) (reference.html#rgb), which is short for “red, green, blue”. RGB is an [additive color model](#) (reference.html#additive-color-model): every shade is some combination of red, green, and blue intensities. We can think of these three values as being the axes in a cube:




Figure: RGB Color Cube

An RGB color is an example of a multi-part value: like a Cartesian coordinate, it is one thing with several parts. We can represent such a value in Python using a [tuple](#) (reference.html#tuple), which we write using parentheses instead of the square brackets used for a list:

```
position = (12.3, 45.6)
print('position is:', position)
color = (10, 20, 30)
print('color is:', color)
```

```
position is: (12.3, 45.6)
color is: (10, 20, 30)
```

We can select elements from tuples using indexing, just as we do with lists and arrays:

```
print('first element of color is:', color[0])
```

```
first element of color is: 10
```

Unlike lists and arrays, though, tuples cannot be changed after they are created — in technical terms, they are [immutable](#) (reference.html#immutable):

```
color[0] = 40
print('first element of color after change:', color[0])
```

```

-----
TypeError                                Traceback (most recent call last)
<ipython-input-11-9c3dd30a4e52> in <module>()
----> 1 color[0] = 40
      2 print('first element of color after change:', color[0])

TypeError: 'tuple' object does not support item assignment

```

If a tuple represents an RGB color, its red, green, and blue components can take on values between 0 and 255. The upper bound may seem odd, but it's the largest number that can be represented in an 8-bit byte (i.e., 2^8-1). This makes it easy for computers to manipulate colors, while providing fine enough gradations to fool most human eyes, most of the time.

Let's see what a few RGB colors actually look like:

```

row = ImageGrid(8, 1)
row[0, 0] = (0, 0, 0) # no color => black
row[1, 0] = (255, 255, 255) # all colors => white
row[2, 0] = (255, 0, 0) # all red
row[3, 0] = (0, 255, 0) # all green
row[4, 0] = (0, 0, 255) # all blue
row[5, 0] = (255, 255, 0) # red and green
row[6, 0] = (255, 0, 255) # red and blue
row[7, 0] = (0, 255, 255) # green and blue
row.show()

```



Simple color values like $(0, 255, 0)$ are easy enough to decipher with a bit of practice, but what color is $(214, 90, 127)$? To help us, `ipythonblocks` provides a function called `show_color`:

```

from ipythonblocks import show_color
show_color(214, 90, 127)

```



It also provides a table of standard colors:

```

from ipythonblocks import colors
c = ImageGrid(3, 2)
c[0, 0] = colors['Fuchsia']
c[0, 1] = colors['Salmon']
c[1, 0] = colors['Orchid']
c[1, 1] = colors['Lavender']
c[2, 0] = colors['LimeGreen']
c[2, 1] = colors['HotPink']
c.show()

```



Making a colorbar

Fill in the ____ in the code below to create a bar that changes color from dark blue to black.

```
bar = ImageGrid(10, 1)
for x in range(10):
    bar[x, 0] = (0, 0, ____)
```

```
bar.show()
```

Why RGB?

Why do computers use red, green, and blue as their primary colors?

Nested loops

Will changing the nesting of the loops in the code above — i.e., wrapping the Y-axis loop around the X-axis loop — change the final image? Why or why not?

Where to change data

Why did we transpose our data outside our heat map function? Why not have the function perform the transpose?

Design choice: return versus display

Why does the heat map function return the grid rather than displaying it immediately? Do you think this is a good or bad design choice?

Software Carpentry (<http://software-carpentry.org>)

Source (<https://github.com/swcarpentry/python-novice-inflammation>)

Contact (<mailto:admin@software-carpentry.org>)

License (<LICENSE.html>)